

Supplementary Material for DeGro: Target Determinacy Verification and Grounded Repair for Executable Mathematical Reasoning

Anonymous Author(s)

This document provides proof details, evaluation protocol, selected extended results, and reproducibility identifiers. The main paper is self-contained.

A Formal Details

Let a structured formalization be $F = (V, D, C, q)$, where V is a finite set of typed variables, D contains domain restrictions, C contains translated source constraints, and $q(V)$ is the target. Write

$$\text{Sol}(F) = \{m \mid m \models D \wedge C\}.$$

The target is determinate when all feasible assignments agree:

$$\text{Det}(F, q) \iff \forall m_1, m_2 \in \text{Sol}(F), q(m_1) = q(m_2). \quad (1)$$

For two disjoint renamings $V^{(1)}$ and $V^{(2)}$, define

$$\Phi(F, q) = F^{(1)} \wedge F^{(2)} \wedge q(V^{(1)}) \neq q(V^{(2)}). \quad (2)$$

Witness equivalence. Assuming semantics-preserving compilation, $\Phi(F, q)$ is satisfiable exactly when two feasible assignments disagree on q . A satisfying assignment projects to the two witnesses; conversely, any such pair satisfies Eq. (2). Hence, if F is satisfiable and $\Phi(F, q)$ is unsatisfiable, Eq. (1) holds. The initial satisfiability check is necessary because an inconsistent F would otherwise yield a vacuous certificate.

Repair gate. For candidate constraint c and a frozen set $\mathcal{R}$ of permitted domain-semantic rules,

$$\begin{aligned} \text{Accept}(c) \iff & \text{Compiles}(c) \wedge \text{Grounded}(c, P, \mathcal{R}) \\ & \wedge \text{SAT}(F \wedge c) \wedge \neg(F \models c) \wedge \text{Det}(F \wedge c, q). \end{aligned} \quad (3)$$

Thus every accepted edit is supported, consistent, novel, and target-sufficient. If neither of two proposals passes, the original F is preserved. Exact-span provenance is an auditable proxy, not a proof of semantic equivalence.

B Reference Algorithm and Trusted Boundary

The bounded controller performs:

1. parse one typed `MODEL SPEC` and reject unsupported expressions;
2. check $\text{SAT}(F)$, then $\Phi(F, q)$;
3. if ambiguous, record two target-divergent witnesses;
4. request one grounded constraint or `ABSTAIN`;
5. apply Eq. (3), stopping after at most two proposals.

The restricted parser, typed compiler, Z3, and deterministic gate form the trusted boundary; the language model does not. The parser rejects undeclared identifiers, arbitrary calls, attribute access, and unsupported operators. Both solving and verification consume the same normalized specification. Logs retain initial and final states, witnesses, proposals, cited spans, rejection reasons, model configuration, and prompt hashes.

C Controlled Paired Evaluation

For a validated gold specification $F^* = \{c_1, \dots, c_k, c^*\}$, define $F^- = F^* \setminus \{c^*\}$. An item is retained only when F^* is satisfiable and target-determinate while F^- is satisfiable and target-ambiguous. Two source conditions share exactly F^- :

- **Omission:** the source retains the sentence supporting c^* ; grounded repair is correct.
- **Underspecified:** that sentence is removed; abstention is correct.

The solver state is identical within each pair, so ambiguity alone cannot reveal the label. The 120-pair confirmatory set is disjoint from the 80-pair pilot and 100-pair development sets and was frozen before outcome inspection. It contains rank-deficient systems (20), linear separation (20), moment problems (40), Chinese-remainder reconstruction (20), and graph path sums (20).

All four policies receive the same source, F^- , target, schema, and two-attempt budget. Self-review receives neither verdict nor grounding rule; grounded self-review adds grounding; determinacy-guided adds the solver verdict; DeGro receives both. Both models use medium reasoning effort and temperature 1.0, giving 960 complete records per model. RSR measures correct repair on omission cases, CAR correct abstention, and FDA = (RSR + CAR)/2. Confidence intervals resample pairs; paired tests use exact two-sided McNemar tests with Holm correction across the two model-specific primary tests.

D Extended Controlled Results

Table 1: Confirmatory results on 120 preregistered pairs per model (percent).

Model	Method	RSR	CAR	FDA	ORR	UR
GPT-OSS 20B	Self-review	93.3	87.5	90.4	12.5	15.2
	Grounded self-review	100.0	80.0	90.0	20.0	16.7
	Determinacy-guided	86.7	99.2	92.9	0.8	6.3
	DeGro	92.5	97.5	95.0	2.5	4.3
Gemma 4 31B	Self-review	100.0	93.3	96.7	6.7	6.3
	Grounded self-review	100.0	95.8	97.9	4.2	4.0
	Determinacy-guided	100.0	99.2	99.6	0.8	0.8
	DeGro	100.0	95.0	97.5	5.0	4.8

For GPT-OSS, DeGro minus grounded self-review is +5.0 FDA points with pair-cluster 95% CI [0.4, 9.6]; exact $p = .050$ and Holm-adjusted $p = .100$. For Gemma, the difference is -0.4 points with CI [-3.3, 2.5] and $p = 1.0$. The evidence therefore does not support a model-independent gain.

Across all 300 pairs, GPT-OSS changes from 91.2% to 96.2% FDA (+5.0, CI [2.3, 7.7], $p < .001$), while Gemma changes from 96.2% to 97.8% (+1.7, CI [-0.2, 3.5], $p = .099$). This cumulative analysis is secondary because it includes development observations. In a 100-pair exploratory ablation, exposing target-divergent witnesses to the repair model did not improve DeGro (98.5% versus 97.5% FDA; adjusted $p = .727$); witnesses remain valuable as audit evidence.

E Natural-Error Study

Exactly one Structured SymCode formalization was generated per problem and reused across methods. Cohort membership was frozen from parse, compiler, and solver states before annotation or repair-outcome inspection. Blinded annotation packets omitted the determinacy verdict, witness, repair, method, and answer outcome.

Two reviewers inspected all 292 MATH-500 and 503 ALG514 final AI-assisted labels. Because separate pre-discussion human labels were not retained, we claim two-person verification, not independent double annotation. Labels are CORRECT, MISSING-CONSTRAINT, WRONG-CONSTRAINT, EXTRA-CONSTRAINT, WRONG-DOMAIN, WRONG-TARGET, COMPUTATIONAL-ERROR, and OTHER; ambiguity itself never determines the label.

Table 2: Human-verified MATH-500 outcomes on 292 conclusive, solver-supported cases (58.4% coverage).

Label	<i>n</i>	Ambiguous	Repair	Solver correct	DeGro correct
Correct	158	33	10	119	126
Missing constraint	24	21	4	2	3
Wrong constraint	32	10	0	16	16
Extra constraint	8	4	1	4	4
Wrong domain	1	1	0	0	0
Wrong target	69	52	10	3	8
Total	292	121	25	144	157

Within the 24 target-critical omissions, ambiguity recall is 87.5% (Wilson 95% CI [69.0, 95.7]), detector precision is 17.4% (CI [11.6, 25.1]), and end-to-end grounded repair is 4.2% (CI [0.7, 20.2]). Overall accuracy changes from 49.3% to 53.8%; 16 discordances favor DeGro and three favor Structured Solver (exact McNemar $p = .00443$).

For ALG514, 503/514 problems are covered (97.86%). Labels comprise 477 correct, 11 wrong-target, eight target-critical missing-constraint, and seven wrong-constraint cases. Structured Solver scores 94.16% and DeGro 94.36% under the released scoring protocol; nine numerically compatible cases still carry a verified semantic error.

F Reproducibility Checklist

Table 3: Frozen artifacts and identifiers.

Artifact	Frozen value
MIRA source revision	72ff918901584d1d1e835a9f86d41e3d51a915a2
MIRA source SHA-256	78a021526f9457af6e3776ac470c4ef1202d18fd9a19539e7e289294997b2ad0
Selection seed	20270911
Confirmatory prompt SHA-256	ea-fed6be4973c21e1efbb5cc804112c9d50c95764e8f5066b58bdee58b92bdda
Confirmatory output SHA-256	5616d0c6203f762560439ee193f1148133d06baf034d6c34c7250a0a44ce3f0b
MATH-500 verified labels	4e1bbe0b9eae-fdefe4569ec10fdd42daf8f1131d17dfbc410d0faed0b430ad38
ALG514 verified labels	72217047ff8d41b4a3766e504325821bd07230d7b6dc7636475294ab8e5ffecd

The released scripts reconstruct tables from JSONL records, preserve cohort provenance, and fail on missing expected counts. Statistical code implements pair-cluster bootstrap intervals, exact McNemar tests, Holm adjustment, and Wilson intervals. All 49 repository tests pass in an isolated environment. Exact prompts, schemas, and per-record logs accompany the code release rather than being reproduced here.

G AI-Assisted Technology Disclosure

OpenAI Codex (GPT-5 model family) assisted with prototype code, experimental scaffolding, language polishing, restructuring, and formatting. It is not an author; the human authors remain responsible for the methodology, implementation, experiments, analyses, citations, and claims. The pipeline image was supplied by the author, who should retain its provenance for venue disclosure if required.